

Final Project Submission

Please fill out:

- Student name: Cynthia Chiuri
- Student pace: part time
- Scheduled project review date/time:
- Instructor name: Noah Kandie
- Blog post URL: <https://datascience942.wordpress.com/2023/06/02/predicting-housing-prices-insights-and-recommendations-for-homeowners/>
(<https://datascience942.wordpress.com/2023/06/02/predicting-housing-prices-insights-and-recommendations-for-homeowners/>)
- GROUP 8



Column Names and Descriptions for King County Data Set

- id - Unique identifier for a house
- date - Date house was sold
- price - Sale price (prediction target)
- bedrooms - Number of bedrooms
- bathrooms - Number of bathrooms
- sqft_living - Square footage of living space in the home
- sqft_lot - Square footage of the lot
- floors - Number of floors (levels) in house
- waterfront - Whether the house is on a waterfront
 - Includes Duwamish, Elliott Bay, Puget Sound, Lake Union, Ship Canal, Lake Washington, Lake Sammamish, other lake, and river/slough waterfronts
- view - Quality of view from house
 - Includes views of Mt. Rainier, Olympics, Cascades, Territorial, Seattle Skyline, Puget Sound, Lake Washington, Lake Sammamish, small lake / river / creek, and other
- condition - How good the overall condition of the house is. Related to maintenance of house.
 - See the [King County Assessor Website](https://info.kingcounty.gov/assessor/esales/Glossary.aspx?type=r) (<https://info.kingcounty.gov/assessor/esales/Glossary.aspx?type=r>) for further explanation of each condition code
- grade - Overall grade of the house. Related to the construction and design of the house.
 - See the [King County Assessor Website](https://info.kingcounty.gov/assessor/esales/Glossary.aspx?type=r) (<https://info.kingcounty.gov/assessor/esales/Glossary.aspx?type=r>) for further explanation of each building grade code
- sqft_above - Square footage of house apart from basement
- sqft_basement - Square footage of the basement

- `yr_built` - Year when house was built
- `yr_renovated` - Year when house was renovated
- `zipcode` - ZIP Code used by the United States Postal Service
- `lat` - Latitude coordinate
- `long` - Longitude coordinate
- `sqft_living15` - The square footage of interior housing living space for the nearest 15 neighbors
- `sqft_lot15` - The square footage of the land lots of the nearest 15 neighbors

Predictive analysis of House prices in King County

Renovations: Worth the Investment or a Risky Gamble?

Overview

This project uses linear regression analysis to infer how certain variables impact housing prices and by how much. The aim is to gain insights and make predictions about the factors that affect house sales in King County area as well as lucrative neighbourhoods to invest in while using statistical techniques to support relevant recommendations.

Business problem

The real estate agency wants to provide homeowners with advice on how home renovations can potentially increase the estimated value of their homes and by what amount. The agency aims to offer valuable insights to homeowners, helping them make informed decisions about renovation projects that can maximize their return on investment when selling their properties.

Business objectives

The analysis aims to answer below questions in trying to predict the prices;

1. To determine how much would adding an extension to the lot area of the home likely increase sale price?
2. To examine how much would adding an additional bathroom likely increase sale price?
3. To determine how much would adding an extension to the living area of the home likely increase sale price?
4. To examine how much would adding an additional floor to a house likely increase sale price?

Metric of Success

Our metric of success will be the R-Squared and the Root Mean Square of Errors(RMSE). This will be the final step in evaluating the performance of the model by doing a train-test split, which will give us an idea of how the model would perform with new data for the same variables that the model will be trained on, and another set that it will be tested on. By default, the function takes 80% of the data as the training subset and the other 20% as its test subset.

Data understanding

The dataset used for predicting the sales price of houses in King County is found in `kc_house_data.csv`. It comprises 21,597 observations and consists of 20 house features along with a column indicating the house price. The data covers homes sold between May 2014 and May 2015. Out of the 20 features, eight are continuous numerical variables that provide information about the area dimensions and geographical location of the house. These variables offer a general overview of the house's structure and characteristics. The remaining attributes are discrete variables, which offer more detailed information about specific components of the house. The discrete variables include quantifications of various items within the house, such as the number of bedrooms, bathrooms, presence of a waterfront, and floor level. Some attributes also provide background information about the house, such as the year of construction, year of innovation, previous selling price, and date of sale.

Importing the relevant libraries and loading the dataset from `kc_house_data.csv`.


```
In [77]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
import statsmodels.formula.api as smf
import statsmodels.stats.api as sms
from statsmodels.compat import lzip
import statsmodels
import math
import matplotlib.pyplot as plt
from scipy.special import logsumexp
from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.datasets import make_regression
from sklearn.linear_model import LinearRegression
import sklearn.metrics as metrics
from scipy import stats as stats
from statsmodels.stats.outliers_influence import variance_inflation_factor
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from statsmodels.formula.api import ols
```

```
In [2]: # displaying the DataFrame
df = pd.read_csv("data/kc_house_data.csv")
df
```

```
Out[2]:
```

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	
...
21592	263000018	5/21/2014	360000.0	3	2.50	1530	1131	3.0	
21593	6600060120	2/23/2015	400000.0	4	2.50	2310	5813	2.0	
21594	1523300141	6/23/2014	402101.0	2	0.75	1020	1350	2.0	
21595	291310100	1/16/2015	400000.0	3	2.50	1600	2388	2.0	
21596	1523300157	10/15/2014	325000.0	2	0.75	1020	1076	2.0	

21597 rows × 21 columns



```
In [3]: #Checking on the columns in our dataset
df.columns
```

```
Out[3]: Index(['id', 'date', 'price', 'bedrooms', 'bathrooms', 'sqft_living',
              'sqft_lot', 'floors', 'waterfront', 'view', 'condition', 'grade',
              'sqft_above', 'sqft_basement', 'yr_built', 'yr_renovated', 'zipcode',
              'lat', 'long', 'sqft_living15', 'sqft_lot15'],
              dtype='object')
```

```
In [4]: # checking the number of rows and columns
df.shape
```

```
Out[4]: (21597, 21)
```

We have 21,597 rows of data, meaning we have information about 21,597 homes. That is plenty of data with which to build a model. However, not every row has complete information about a given home, such as `yr_renovated` having fewer than 21,597 records.

```
In [5]: # checking the summary statistics
df.describe()
```

```
Out[5]:
```

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot
count	2.159700e+04	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04
mean	4.580474e+09	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04
std	2.876736e+09	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04
min	1.000102e+06	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02
25%	2.123049e+09	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06

This gives us a great overview of the data we have. A few key takeaways are:

- Homes are priced between 78,000 and 7,700,000 dollars
- Most homes are between 322,000 and 645,000 dollars
- The average home has 3.3 bedrooms and 2.1 bathrooms, with about 2,080 living square footage
- All homes have between 1 and 3.5 floors
- The average home was built around 1971, but some are over 100 years old
- We noticed that there is a home listed as having 33 bedrooms. Either that's an extreme outlier, or some sort of input error. We will investigate that later.

So now that we have a basic understanding of the data we're working with, we can dive into some more information that we will need in order to build a model later. By using the `.info()` method, we can pull up information about missing data values, how many rows of data we have, and whether values are being read as text or as numerical data.

```
In [6]: # checking the metadata of our data
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   id                     21597 non-null  int64
1   date                   21597 non-null  object
2   price                  21597 non-null  float64
3   bedrooms               21597 non-null  int64
4   bathrooms              21597 non-null  float64
5   sqft_living            21597 non-null  int64
6   sqft_lot               21597 non-null  int64
7   floors                 21597 non-null  float64
8   waterfront             19221 non-null  object
9   view                   21534 non-null  object
10  condition              21597 non-null  object
11  grade                  21597 non-null  object
12  sqft_above             21597 non-null  int64
13  sqft_basement          21597 non-null  object
14  yr_built               21597 non-null  int64
15  yr_renovated           17755 non-null  float64
16  zipcode                21597 non-null  int64
17  lat                    21597 non-null  float64
18  long                   21597 non-null  float64
19  sqft_living15          21597 non-null  int64
20  sqft_lot15             21597 non-null  int64
dtypes: float64(6), int64(9), object(6)
memory usage: 3.5+ MB
```

From the metadata, not every row has complete information about a given home, such as `yr_renovated` having fewer than 21,597 entries.

Furthermore, not all columns of data are being read as quantitative data. In this case, some columns are being read as numbers, whether that's in integer form or float (numbers with decimals) form, while others are being read as text inputs, or objects.

It looks like we'll have to convert some columns with qualitative data (such as `view`, `waterfront`, and `condition`) into integers or floats so we can build models with them. We'll also have to replace null values for the `waterfront`, `view`, and `yr_renovated` columns.

```
In [7]: # checking for the total number of null values per column
df.isna().sum()
```

```
Out[7]: id                0
        date              0
        price             0
        bedrooms          0
        bathrooms         0
        sqft_living       0
        sqft_lot          0
        floors            0
        waterfront       2376
        view              63
        condition        0
        grade            0
        sqft_above        0
        sqft_basement     0
        yr_built          0
        yr_renovated     3842
        zipcode           0
        lat               0
        long              0
        sqft_living15     0
        sqft_lot15        0
        dtype: int64
```

Based on the dataset waterfront, view and yr_renovated have the summation of 2,376, 63 and 3,842 null values respectively.

```
In [8]: # dropping null values
df.dropna(inplace=True)
```



```
In [9]: # checking if the null values are successfully dropped.  
df.isna().sum()
```

```
Out[9]: id                0  
date                0  
price              0  
bedrooms           0  
bathrooms          0  
sqft_living        0  
sqft_lot           0  
floors             0  
waterfront         0  
view               0  
condition          0  
grade              0  
sqft_above         0  
sqft_basement      0  
yr_built           0  
yr_renovated       0  
zipcode            0  
lat                0  
long               0  
sqft_living15      0  
sqft_lot15         0  
dtype: int64
```

```
In [10]: # checking on duplicated values in id column.  
duplicated=df["id"].duplicated().sum()  
duplicated
```

```
Out[10]: 86
```

This shows that there are 86 duplicates in the `id` column. This is equivalent to 86 houses from the the dataset. Dropping the mentioned number may not skew the dataset.

```
In [11]: # dropping the duplicates  
df.drop_duplicates(subset='id', inplace=True)
```

```
In [12]: # confirming that the duplicates have been dropped successfully  
duplicated=df["id"].duplicated().sum()  
duplicated
```

```
Out[12]: 0
```

Exploratory Data Analysis (EDA)

Univariate analysis

The stage involve exploration process, which involves generating and plotting histograms and box plots. This crucial step allows us to gain insight into the distribution patterns of the data for each variable. By visualizing the histograms, we can better comprehend the spread and frequency of values within each variable, providing a foundation for further analysis. Box plots help us identify potential outliers.

In [13]: *# Checking on measures of central tendency and dispersion*


```
price_mean = df["price"].mean()
price_mode = df["price"].mode()[0]
price_median = df["price"].median()
price_std = df["price"].std()

print("Mean:", price_mean)
print("Mode:", price_mode)
print("Median:", price_median)
print("Standard Deviation:", price_std)
```

Mean: 541492.6832737944
Mode: 350000.0
Median: 450000.0
Standard Deviation: 372603.68455896684

In [14]: *# Plotting a histogram of price*


```
plt.hist(df["price"], bins=10, edgecolor='black')
plt.xlabel("Price")
plt.ylabel("Frequency")
plt.title("Histogram of Prices")
plt.show()
```



```
In [15]: # Plotting a histogram/kernel density estimate of price

plt.figure(figsize=(10,4))


price_dist = sns.distplot(df["price"])
price_dist.set(xlabel="Price in Millions", title="Price Density of Houses in K")
plt.show()
```



As we can see, the distribution of house prices is right-skewed. This means that there are a large number of houses that are relatively inexpensive, but there are also a small number of houses that are very expensive.

```
In [16]: # Checking on outliers in the price variable

plt.figure()
plt.boxplot(df['price'])
plt.ylabel("feature")
plt.title('Box Plot of Price')
plt.show()
```



Based on the box plot there is presence of outliers but we decided to keep them based on the assumption that they are a true representation of the real-world dataset.

```

In [17]: # Plotting Histogram, density plots and box plot

# Select the desired features
features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'zipcode']
plt.figure(figsize=(12, 8))
ncols=3
nrows=4
# Perform univariate analysis for each feature
for feature in features:
    # Descriptive Statistics
    print('Descriptive Statistics for', feature)
    print(df[feature].describe())
    print()
    # Histogram
    plt.figure()
    plt.hist(df[feature], bins=20)
    plt.xlabel(feature)
    plt.ylabel('Frequency')
    plt.title('Histogram of ' + feature)
    plt.show()
    # Box Plot
    plt.figure()
    plt.boxplot(df[feature])
    plt.ylabel(feature)
    plt.title('Box Plot of ' + feature)
    plt.show()
    # Density Plot
    plt.figure()
    df[feature].plot(kind='density')
    plt.xlabel(feature)
    plt.ylabel('Density')
    plt.title('Density Plot of ' + feature)
    plt.show()

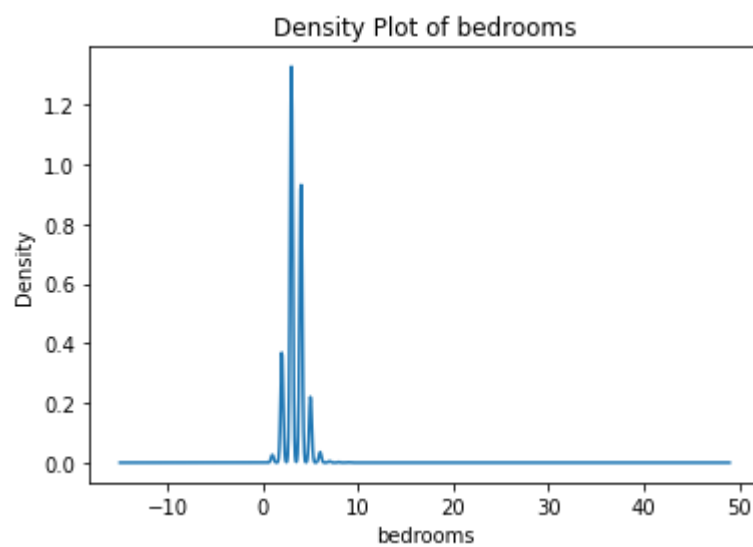


```

Descriptive Statistics for bedrooms

count	15676.000000
mean	3.379434
std	0.935193
min	1.000000
25%	3.000000
50%	3.000000
75%	4.000000
max	33.000000

Name: bedrooms, dtype: float64


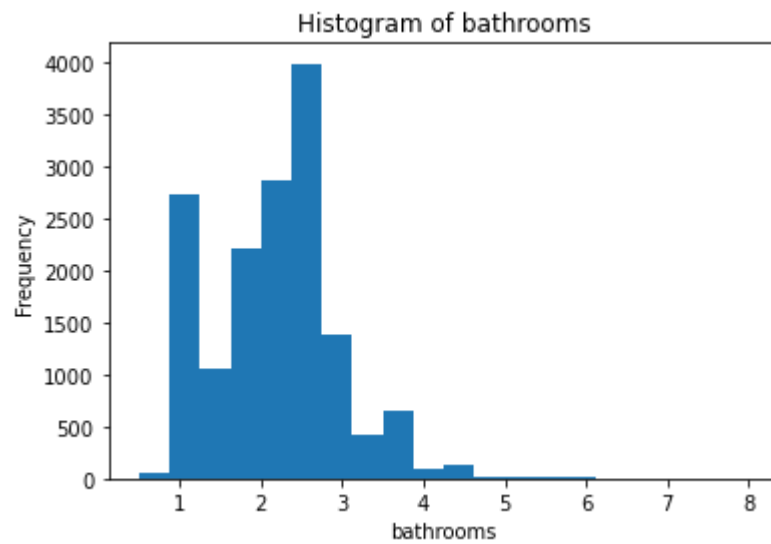
<Figure size 864x576 with 0 Axes>




Descriptive Statistics for bathrooms

```
count    15676.000000
mean       2.122066
std        0.766735
min        0.500000
25%        1.750000
50%        2.250000
75%        2.500000
max        8.000000
```

Name: bathrooms, dtype: float64





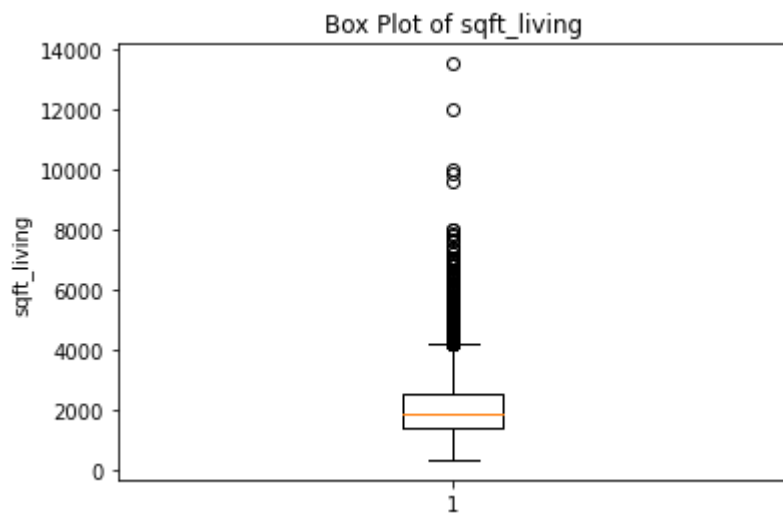


Descriptive Statistics for sqft_living

count	15676.000000
mean	2086.057285
std	918.753332
min	370.000000
25%	1430.000000
50%	1920.000000
75%	2550.000000
max	13540.000000

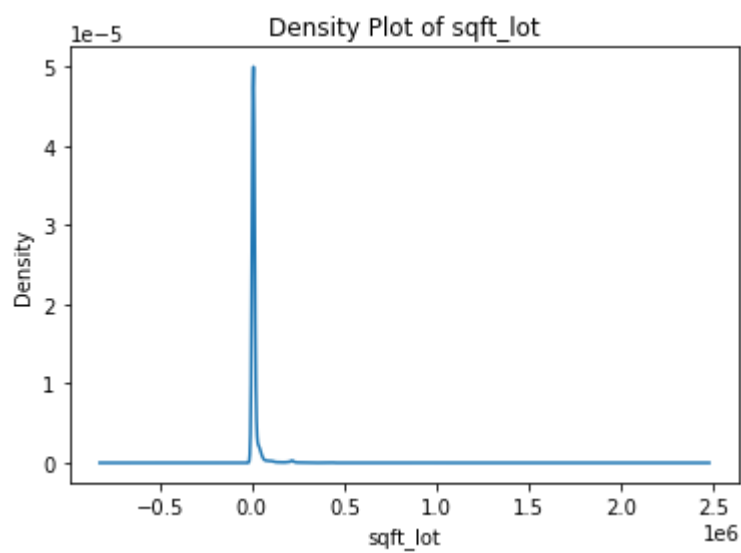

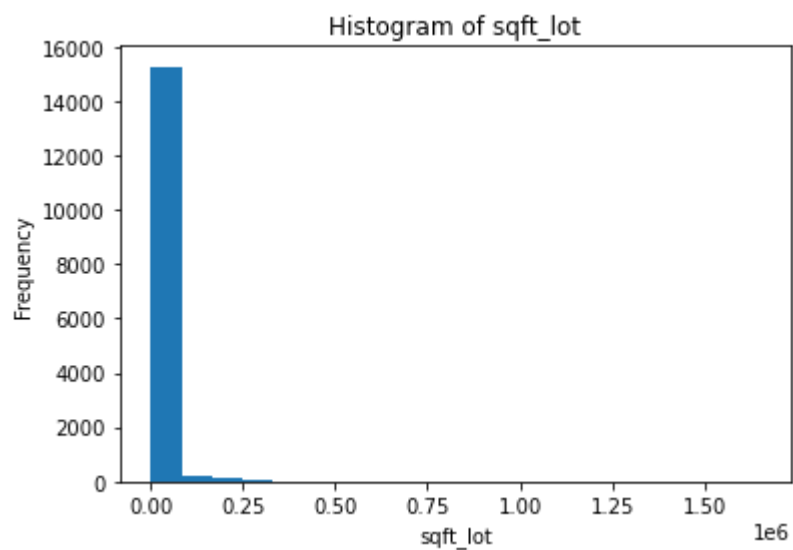
Name: sqft_living, dtype: float64





Descriptive Statistics for sqft_lot



```
count    1.567600e+04
mean      1.529400e+04
std       4.189635e+04
min       5.200000e+02
25%       5.045250e+03
50%       7.600000e+03
75%      1.071700e+04
max      1.651359e+06
Name: sqft_lot, dtype: float64
```




Descriptive Statistics for floors

```
count    15676.000000  
mean      1.496587  
std       0.539689  
min       1.000000  
25%      1.000000  
50%      1.500000  
75%      2.000000  
max       3.500000
```


```
Name: floors, dtype: float64
```


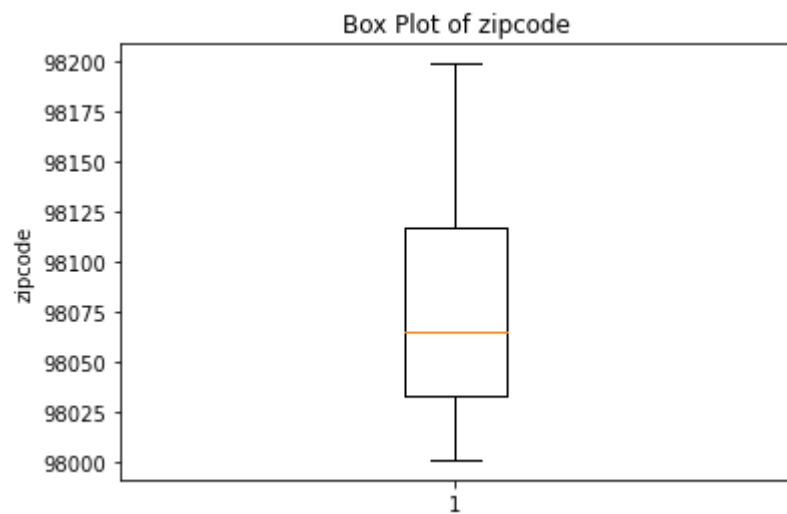




Descriptive Statistics for zipcode

```
count    15676.000000
mean     98077.487114
std       53.366170
min      98001.000000
25%      98033.000000
50%      98065.000000
75%      98117.000000
max      98199.000000
Name: zipcode, dtype: float64
```





Descriptive Statistics for sqft_living15

count 15676.000000

mean 1991.289168

std 684.179299

min 399.000000




25% 1490.000000

50% 1850.000000

75% 2370.000000

max 6210.000000


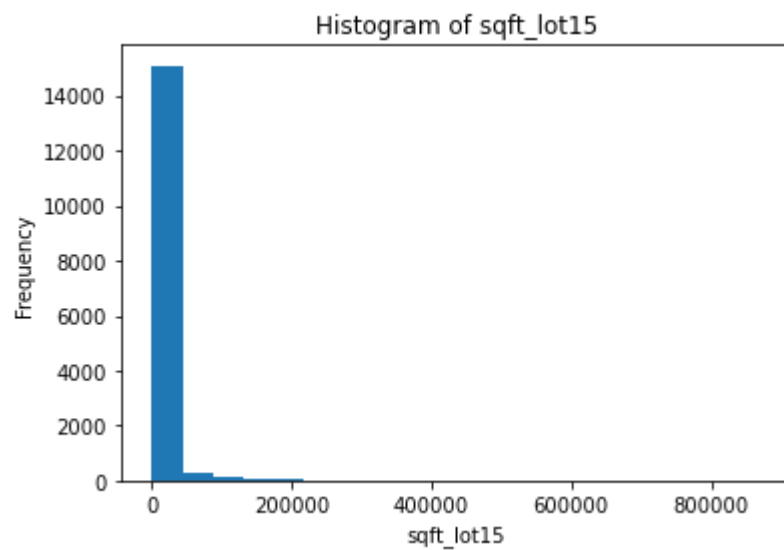
Name: sqft_living15, dtype: float64

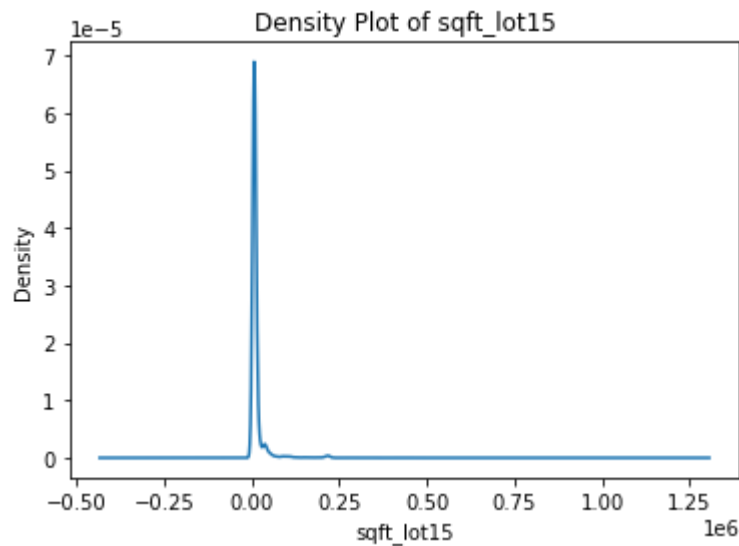


Descriptive Statistics for sqft_lot15

```
count    15676.000000
mean     12911.040125
std      28037.170327
min       659.000000
25%      5100.000000
50%      7620.000000
75%     10102.250000
max     871200.000000
```

```
Name: sqft_lot15, dtype: float64
```





Descriptive Statistics

In [18]: *# To ascertain the median points of the dataset*

```
df[['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors', 'sqft_living15', 'sqft_lot15']]
```

```
Out[18]: bedrooms      3.00
bathrooms      2.25
sqft_living    1920.00
sqft_lot      7600.00
floors         1.50
sqft_living15  1850.00
sqft_lot15    7620.00
dtype: float64
```

- Based on the above plots, bedrooms shows out of 15,676 counts the mean mean is 3.37, std of 0.935 and the median of 3.0. This shows most of the houses have 3 or 4 bedrooms with an exception of an outlier which exist in the dataset. The dataset is also distributed uniformly around the mean.
- Bathrooms has a mean of 2.12, median of 2.25 and std of 0.77. The dataset is rightly skewed and the dataset is distributed around the mean. Most of the houses have 2 bathrooms.
- Sqft_living shows that it has a mean of 2080, std of 918.1 and median of 1910 depicting that the data is distributed around the mean with slight deviations. most of the houses covers 2080 square feet space.
- Sqft_lot has rightly skewed dataset with the presence of outlier. It has a mean of 15,099.41 feet, median of 7,618 which shows that few data points are around the mean.
- Sqft_living15 dataset shows that the dataset has mean of 1,986 feats of living space, median of 1,840 which shows that most of the houses have living space of 1,986 feet and since most of the datapoints are distributed around the mean with a deviation of 685 feet only.
- Descriptive Statistics for floors depicts that relatively few houses have 1 to 2 foors. The mean is 1.5, median of 1.5. the data points are scattered with most points a 1 and 2 based on the density curve.

Bivariate Analysis

Bivariate analysis focuses on determining the correlation between two variables. At this stage we will use a heatmap and scatterplot to check on correlation and collinearity of the variables.

```
In [19]: features = ['bedrooms', 'bathrooms', 'sqft_living', 'sqft_lot', 'floors',
                    'zipcode', 'sqft_living15', 'sqft_lot15', 'yr_built']

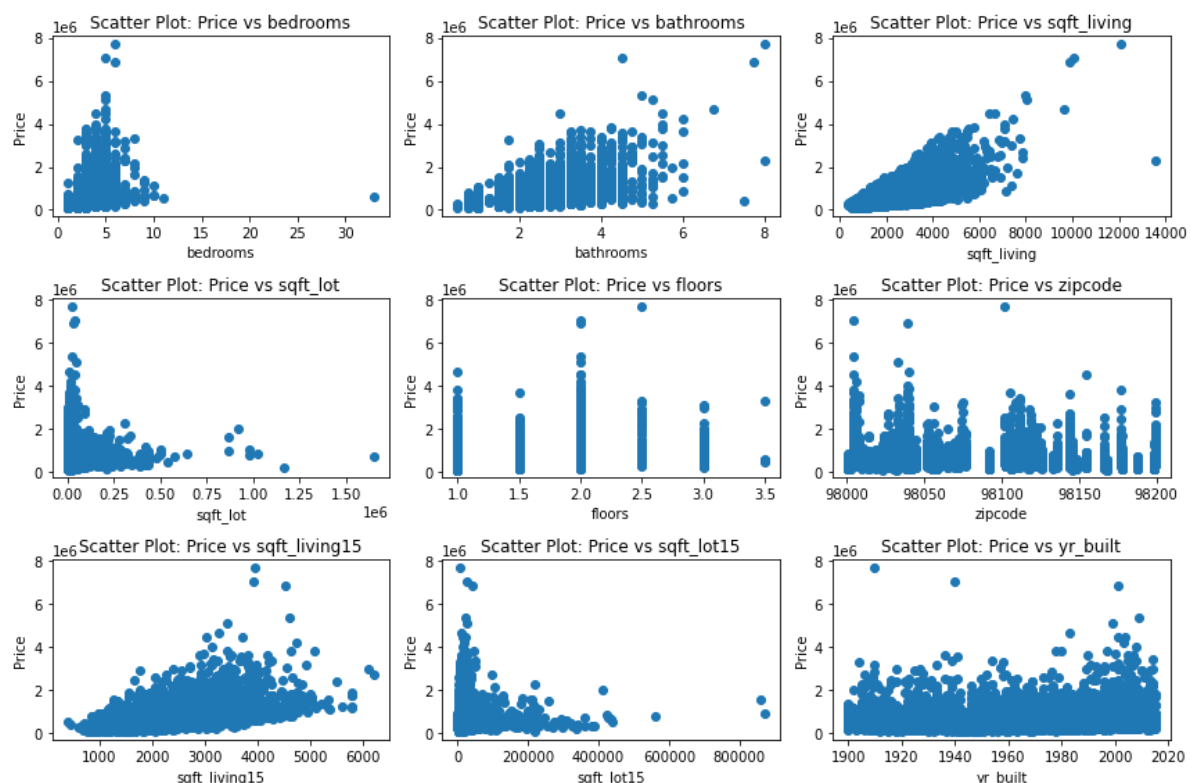
# Set the figure size and grid layout
fig, axs = plt.subplots(nrows=3, ncols=3, figsize=(12, 8))

# Perform bivariate analysis for each feature
for i, feature in enumerate(features):
    # Calculate the row and column index
    row = i // 3
    col = i % 3

    # Scatter Plot
    axs[row, col].scatter(df[feature], df['price'])
    axs[row, col].set_xlabel(feature)
    axs[row, col].set_ylabel('Price')
    axs[row, col].set_title('Scatter Plot: Price vs ' + feature)

# Adjust the spacing between subplots
plt.tight_layout()

# Show the plot
plt.show()
```



- Square foot of living has a STRONG correlation with price; we can assume that as the square foot of living increases, so does price.
- Square foot of lot has a high number of 0's. What does this mean? Does this indicate apartment building homes, which is more expansive vertically rather than horizontally (compared to regular flat homes), thus requiring not that much square foot of lot.

In [20]: *# Selecting the numeric columns*

```
numeric_columns = df.select_dtypes(include='number').columns
numeric_df = df[numeric_columns]
```

In [21]: *# Computing the correlation matrix to check for Linearity*

```
numeric_df.corr()
```


Out[21]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_
id	1.000000	-0.016236	0.007883	0.005406	-0.008858	-0.136009	0.020083	-0.0
price	-0.016236	1.000000	0.305947	0.526228	0.705975	0.083572	0.259193	0.6
bedrooms	0.007883	0.305947	1.000000	0.512488	0.574179	0.025684	0.180158	0.4
bathrooms	0.005406	0.526228	0.512488	1.000000	0.753613	0.080027	0.504916	0.6
sqft_living	-0.008858	0.705975	0.574179	0.753613	1.000000	0.164512	0.358657	0.8
sqft_lot	-0.136009	0.083572	0.025684	0.080027	0.164512	1.000000	-0.010454	0.1
floors	0.020083	0.259193	0.180158	0.504916	0.358657	-0.010454	1.000000	0.5
sqft_above	-0.009551	0.611886	0.474835	0.685456	0.876260	0.173422	0.528179	1.0
yr_built	0.024011	0.048672	0.153048	0.504193	0.313206	0.051256	0.486854	0.4
yr_renovated	-0.010419	0.123077	0.016632	0.047255	0.049992	0.002169	-0.001287	0.0
zipcode	-0.007812	-0.048661	-0.148417	-0.198798	-0.195836	-0.129495	-0.057011	-0.2
lat	-0.006173	0.306058	-0.007583	0.029184	0.057228	-0.084771	0.058032	0.0
long	0.018679	0.020241	0.129424	0.221825	0.238786	0.231748	0.128729	0.3
sqft_living15	0.000362	0.580963	0.392272	0.569053	0.756576	0.144640	0.281330	0.7
sqft_lot15	-0.141551	0.078972	0.025342	0.081837	0.176506	0.718327	-0.013882	0.1

In [22]: *# Creating a heatmap using seaborn*

```
columns=['price', 'bedrooms', 'grade_no', 'yr_built', 'sqft_living', 'floors',
        'bathrooms', 'cond_avg', 'cond_fair', 'cond_good', 'cond_poor', 'cond_v']
index=['price', 'bedrooms', 'grade_no', 'yr_built', 'sqft_living', 'floors',
        'bathrooms', 'cond_avg', 'cond_fair', 'cond_good', 'cond_poor', 'cond_v']

corr_matrix = numeric_df.corr()
mask = np.triu(np.ones_like(corr_matrix, dtype=bool))
fig, ax = plt.subplots(figsize=(10, 8))
sns.heatmap(corr_matrix, mask=mask, annot=True, cmap='coolwarm', fmt=".2f", li
ax.set_title('Correlation Heatmap')
plt.xticks(rotation=90)
plt.yticks(rotation=0)
plt.show()
```



Data Pre-processing before fitting our Regression Model

This involves techniques such as:

1. Deal with null values
2. Encoding categorical variables

3. Feature engineering
4. Transformations
5. Feature scaling

```
In [23]: # converting sqft_basement and waterfront which involves using OneHotEncoder.
df['sqft_basement'] = pd.to_numeric(df['sqft_basement'], errors='coerce')
```

Categorical columns include condition and waterfront .

One Hot Encoding the Categorical Variables

```
In [24]: # One_Hot_Encoding the categorical variables

df["grade_no"] = pd.to_numeric(df['grade'].str.split().str[0])

condition = df[['condition']]
ohe = OneHotEncoder(categories="auto", sparse=False, handle_unknown="ignore")
ohe.fit(condition)
condition_enc = ohe.transform(condition)
condition_enc = pd.DataFrame(condition_enc,
                             columns=['cond_avg', 'cond_fair', 'cond_good', 'cond_
                                     index=df.index)
df.drop('condition', axis=1, inplace=True)
df = pd.concat([df, condition_enc], axis=1)
```

```
In [25]: # Selecting our features of relevance

df_values = df[['price', 'bedrooms', 'bathrooms', 'sqft_living', 'sqft_basement',
                'sqft_lot15', 'grade_no', 'cond_avg', 'cond_fair', 'cond_good',
                'cond_poor', 'cond_verygood']]
```

In [26]: *# Confirming if there are any null values*

```
df.isna().sum()
```

Out[26]:

id	0
date	0
price	0
bedrooms	0
bathrooms	0
sqft_living	0
sqft_lot	0
floors	0
waterfront	0
view	0
grade	0
sqft_above	0
sqft_basement	332
yr_built	0
yr_renovated	0
zipcode	0
lat	0
long	0
sqft_living15	0
sqft_lot15	0
grade_no	0
cond_avg	0
cond_fair	0
cond_good	0
cond_poor	0
cond_verygood	0
dtype:	int64

In [27]: *# Replacing the the null values with 0*

```
df['sqft_basement'] = df['sqft_basement'].fillna(0)
```

In [28]: *# Checking if the null values have been replaced with 0*

```
df.isna().sum()
```

Out[28]:

id	0
date	0
price	0
bedrooms	0
bathrooms	0
sqft_living	0
sqft_lot	0
floors	0
waterfront	0
view	0
grade	0
sqft_above	0
sqft_basement	0
yr_built	0
yr_renovated	0
zipcode	0
lat	0
long	0
sqft_living15	0
sqft_lot15	0
grade_no	0
cond_avg	0
cond_fair	0
cond_good	0
cond_poor	0
cond_verygood	0
dtype:	int64


In [29]: *# Displaying our final df before modeling*

df_values

Out[29]:

	price	bedrooms	bathrooms	sqft_living	sqft_basement	sqft_lot15	grade_no	cond_
1	538000.0	3	2.25	2570	400.0	7639	7	
3	604000.0	4	3.00	1960	910.0	5000	7	
4	510000.0	3	2.00	1680	0.0	7503	8	
5	1230000.0	4	4.50	5420	1530.0	101930	11	
6	257500.0	3	2.25	1715	NaN	6819	7	
...
21591	475000.0	3	2.50	1310	130.0	1265	8	
21592	360000.0	3	2.50	1530	0.0	1509	8	
21593	400000.0	4	2.50	2310	0.0	7200	8	
21594	402101.0	2	0.75	1020	0.0	2007	7	
21596	325000.0	2	0.75	1020	0.0	1357	7	

15676 rows × 12 columns



LINEAR MODELING

Checking for the Linearity Assumption.

Here, we assert two things before building our model;

1. We want to include the features which have the highest correlation with our target variable(price).
2. While following the condition above, we want our features not to be multicorrelated with each other.

In [30]: *# checking for correlations between our features and the target variable
from the highest to the lowest*

```
df.corr()['price'].sort_values(ascending=False).head(15)
```

Out[30]:

price	1.000000
sqft_living	0.705975
grade_no	0.664092
sqft_above	0.611886
sqft_living15	0.580963
bathrooms	0.526228
sqft_basement	0.315663
lat	0.306058
bedrooms	0.305947
floors	0.259193
yr_renovated	0.123077
sqft_lot	0.083572
sqft_lot15	0.078972
cond_verygood	0.055422
yr_built	0.048672

Name: price, dtype: float64

In [31]: *# Checking for Multicollinearity in our predictors*
corr_df = df.corr().abs().stack().reset_index().sort_values(0, ascending=False)
corr_df['pairs'] = list(zip(corr_df.level_0, corr_df.level_1))

Dropping 'level_0' and 'level_1'
corr_df.set_index(['pairs'], inplace=True)
corr_df.drop(columns=['level_0', 'level_1'], inplace=True)

Renaming our column
corr_df.columns = ["corr_coef"]

*# Viewing the highly correlated predictor pairs
(our threshold is features with a value above 80%)*

```
corr_df[(corr_df.corr_coef > 0.80) & (corr_df.corr_coef < 1)]
```

Out[31]:

	corr_coef
pairs	
(sqft_living, sqft_above)	0.876260
(sqft_above, sqft_living)	0.876260
(cond_avg, cond_good)	0.811063
(cond_good, cond_avg)	0.811063

In [32]: *# Dropping unnecessary columns*
df.drop(columns=['id', 'date', 'grade', 'yr_built', 'yr_renovated', 'lat', 'long',
'cond_fair', 'cond_good', 'cond_poor', 'cond_verygood'], inplace=


```
In [33]: # Checking the metadata of the remaining columns
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 15676 entries, 1 to 21596
Data columns (total 14 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   price               15676 non-null  float64
1   bedrooms            15676 non-null  int64
2   bathrooms           15676 non-null  float64
3   sqft_living         15676 non-null  int64
4   sqft_lot            15676 non-null  int64
5   floors              15676 non-null  float64
6   waterfront          15676 non-null  object
7   view                15676 non-null  object
8   sqft_above          15676 non-null  int64
9   sqft_basement       15676 non-null  float64
10  zipcode              15676 non-null  int64
11  sqft_living15        15676 non-null  int64
12  sqft_lot15           15676 non-null  int64
13  grade_no             15676 non-null  int64
dtypes: float64(4), int64(8), object(2)
memory usage: 1.8+ MB
```

Defining our Functions for use

```
In [34]: # Defining a function for fitting our model
def run_model(data):
    x = data.drop('price', axis=1)
    y = data['price']
    linreg = LinearRegression()
    crossvalidation = KFold(n_splits = 10, shuffle = True, random_state = 1)
    mean_r2 = np.mean(cross_val_score(linreg, x, y, scoring='r2', cv=crossvalidation))
    mse = np.mean(cross_val_score(linreg, x, y, scoring='neg_mean_squared_error', cv=crossvalidation))
    rmse = np.sqrt(mse)

    x_cols = data.drop('price', axis=1).columns
    y_col = 'price'
    plus = '+'.join(x_cols)
    formula = y_col + '~' + plus
    model = ols(formula=formula, data=data).fit()
    print('The mean r^2 for a KFold test with 10 splits is {} \n'.format(mean_r2))
    print('The mean RMSE for a KFold test with 10 splits is {} \n'.format(rmse))
    print(model.summary())

    # Testing for homoscedasticity
    residuals = model.resid
    fig, ax = plt.subplots(figsize=(15,8))
    plt.scatter(model.predict(x), residuals)
    plt.plot(model.predict(x), [np.mean(residuals) for i in range(len(data))])
    ax.set_title('Homoscedasticity')
    plt.show()
    print('\n')

    # Testing for normality using a QQ-plot
    fig, ax = plt.subplots(figsize=(15,8))
    sm.graphics.qqplot(residuals, dist=stats.norm, line='45', fit=True, ax=ax)
    ax.set_title('QQ Plot')
    plt.show()
```

```
In [35]: # Defining a function to perform log transformations
def log_transform(features, df):
    for feat in features:
        df[feat] = df[feat].map(lambda x: np.log(x))
    return df
```

```
In [36]: # Defining a function to generate a heatmap
def heatmap(data):
    corr = data.corr()
    fig, ax = plt.subplots(figsize=(12,12))
    sns.heatmap(corr, cmap='Reds', annot=True, ax=ax);
```

```

In [37]: # Defining a function to remove outliers from our features
def outliers(features, data):
    for feat in features:
        mu = np.mean(data[feat])
        std = np.std(data[feat])
        outlier = 3*std
        data = data[(data[feat] <= mu+outlier) & (data[feat] >= mu-outlier)]
    return data

In [38]: # Defining a function to perform OneHotEncoding
def scale_ohe(ohe_feature, data):
    ohe = pd.get_dummies(data[ohe_feature], prefix=ohe_feature, drop_first=True)
    no_ohe = data.drop(ohe_feature, axis=1)
    no_ohe_scale = no_ohe.apply(scale)
    return pd.concat([no_ohe_scale, ohe], axis=1)

In [39]: # Defining a function for getting the coefficients of features
def get_coefficients_continuous(scaled_coefs, features):
    for i, feat in enumerate(features):
        maximum = df_log['price'].max()
        minimum = df_log['price'].min()
        range_feat = df_no_outlier[feat].max() - df_no_outlier[feat].min()
        unscale = abs(scaled_coefs[i])*(maximum-minimum)+minimum
        unlog = math.exp(unscale)

        slope_actual = unlog/range_feat

        if scaled_coefs[i] >= 0:
            print('Coefficient for {} is {}'.format(feat, slope_actual))
        else:
            print('Coefficient for {} is {}'.format(feat, slope_actual*-1))

In [40]: df = df[df['sqft_basement'] != '?']
df['sqft_basement'] = df['sqft_basement'].astype(float)

In [41]: df['sqft_basement'] = df['sqft_basement'].astype(float)

In [42]: df['basement'] = np.where(df['sqft_basement'] > 0, 1, 0)
df.head()

```

Out[42]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	sqft_above
1	538000.0	3	2.25	2570	7242	2.0	NO	NONE	2170
3	604000.0	4	3.00	1960	5000	1.0	NO	NONE	1050
4	510000.0	3	2.00	1680	8080	1.0	NO	NONE	1680
5	1230000.0	4	4.50	5420	101930	1.0	NO	NONE	3890
6	257500.0	3	2.25	1715	6819	2.0	NO	NONE	1715

```
In [43]: df.drop(columns=["sqft_basement", "waterfront", "view"], inplace=True)
```

Building the Baseline model

For the baseline model, we will do a simple linear regression, using the most highly correlated feature and then we improve our model from there through an iterative process whereby we perform techniques such as:

1. Dealing with outliers, i.e. either removing outliers or apply transformations to make the data more robust to outliers.
2. Transformations e.g. log transformations of our features.
3. Feature Scaling, i.e. to ensure that all features are on a similar scale. Common scaling techniques include standardization (mean normalization) or normalization (min-max scaling). This will aid in direct comparison of our features and determine which has the highest impact on our target variable.

```
In [44]: # Assigning our features and target variables
X = df["sqft_living"]
y = df['price']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random

# Baseline Model with statsmodels
X_train_with_intercept = sm.add_constant(X_train)
baseline_model = sm.OLS(y_train, X_train_with_intercept)
baseline_results = baseline_model.fit()
baseline_predictions = baseline_results.predict(sm.add_constant(X_test))
baseline_rmse = mean_squared_error(y_test, baseline_predictions, squared=False

print("Baseline Model RMSE:", baseline_rmse)
print(baseline_results.summary())
```

Baseline Model RMSE: 271201.25051764137

OLS Regression Results

```

=====
=
Dep. Variable:          price    R-squared:                0.49
5
Model:                  OLS      Adj. R-squared:            0.49
5
Method:                 Least Squares    F-statistic:            1.229e+0
4
Date:                   Thu, 01 Jun 2023    Prob (F-statistic):      0.0
0
Time:                   12:26:49    Log-Likelihood:          -1.7425e+0
5
No. Observations:       12540    AIC:                     3.485e+0
5
Df Residuals:           12538    BIC:                     3.485e+0
5
Df Model:                1
Covariance Type:        nonrobust
=====

```

```

=====
==
               coef      std err          t      P>|t|      [0.025      0.97
5]
-----
--
const      -5.235e+04    5849.737      -8.949      0.000     -6.38e+04     -4.09e+
04
sqft_living    285.1177      2.572     110.841      0.000      280.076      290.1
60
=====

```

```

=====
=
Omnibus:                8675.250    Durbin-Watson:            2.00
5
Prob(Omnibus):           0.000    Jarque-Bera (JB):          331404.03
7
Skew:                    2.839    Prob(JB):                  0.0
0
Kurtosis:                27.536    Cond. No.                  5.68e+0
3
=====
=

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The condition number is large, 5.68e+03. This might indicate that there are strong multicollinearity or other numerical problems.

Interpretation of results

1. The model is generally statistically significant with an F-statistic p_value of 0.0 at a significance level of 0.05

2. The R-squared value is 0.495, indicating that approximately 49.5% of the variation in the price can be explained by the sqft_living variable. This value is very low and the model needs improving.
3. The coefficient of the constant term (const) is -5.235×10^4 , and the coefficient of the sqft_living variable is 285.1177. These coefficients represent the expected change in the price for a one-unit change in the corresponding predictor variable, assuming other variables are held constant, e.g. For a one-unit increase in square-foot living area, we see an associated increase in around 285 dollars in selling price of the houses.

Iteration 1

Here we perform the first iteration whereby we have included more features into the model. We also perform a KFold test with 10 splits and get the mean r-squared as well as the mean RMSE of our model.

In [45]: *# Fit our model using the defined function*

```
run_model(df)
```


The mean r^2 for a KFold test with 10 splits is 0.558608548218698

The mean RMSE for a KFold test with 10 splits is 247521.86276668686

OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.56
1
Model:                  OLS      Adj. R-squared:            0.56
1
Method:                 Least Squares    F-statistic:          181
9.
Date:                   Thu, 01 Jun 2023    Prob (F-statistic):      0.0
0
Time:                   12:26:50    Log-Likelihood:          -2.1689e+0
5
No. Observations:       15676    AIC:                    4.338e+0
5
Df Residuals:           15664    BIC:                    4.339e+0
5
Df Model:               11
Covariance Type:        nonrobust
=====
```


```
=====
=====
coef      std err          t      P>|t|      [0.025      0.
975]
-----
----
Intercept    -6.288e+07    3.92e+06    -16.060    0.000    -7.06e+07    -5.52
e+07
bedrooms     -4.35e+04    2676.343    -16.255    0.000    -4.88e+04    -3.83
e+04
bathrooms    -1.275e+04    4454.806     -2.863    0.004    -2.15e+04    -402
1.685
sqft_living   255.7224      8.380     30.514    0.000     239.296     27
2.149
sqft_lot       0.0359      0.068      0.529    0.597     -0.097
0.169
floors        -2.511e+04    5016.612     -5.005    0.000    -3.49e+04    -1.53
e+04
sqft_above    -40.1904      9.273     -4.334    0.000     -58.366     -2
2.015
zipcode       635.9363     39.909     15.935    0.000     557.710     71
4.162
sqft_living15  26.4631      4.864      5.440    0.000     16.928      3
5.998
sqft_lot15    -0.6038      0.102     -5.904    0.000     -0.804      -
0.403
grade_no      9.843e+04    3008.590     32.717    0.000     9.25e+04     1.04
e+05
basement      5038.0897    7251.047      0.695    0.487    -9174.799     1.93
e+04
=====
```


```
=====
Omnibus:           12544.377    Durbin-Watson:          1.97
6
```

Prob(Omnibus):	0.000	Jarque-Bera (JB):	823131.54
1			
Skew:	3.365	Prob(JB):	0.0
0			
Kurtosis:	37.856	Cond. No.	2.00e+08
8			
=====			
=			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 2e+08. This might indicate that there are strong multicollinearity or other numerical problems.






Interpretation of results

1. The model is generally statistically significant with an F-statistic p_value of 0.0 at a significance level of 0.05
2. The R-squared value is 0.561, indicating that approximately 56.1% of the variation in the price can be explained by the model. This value indicates an improvement of the baseline model.
3. The plot to test for homoscedasticity reveals that the residuals are somewhat heteroscedastic because they are diverging/varying. This is an indication of skewness/heavy-tailed dataset/presence of outliers.
4. The QQ-plot is used to test for normality of residuals. In this case, the residuals appear not to be normal because they are diverging off the line.

```
In [46]: # Generating the heatmap
heatmap(df)
```



```

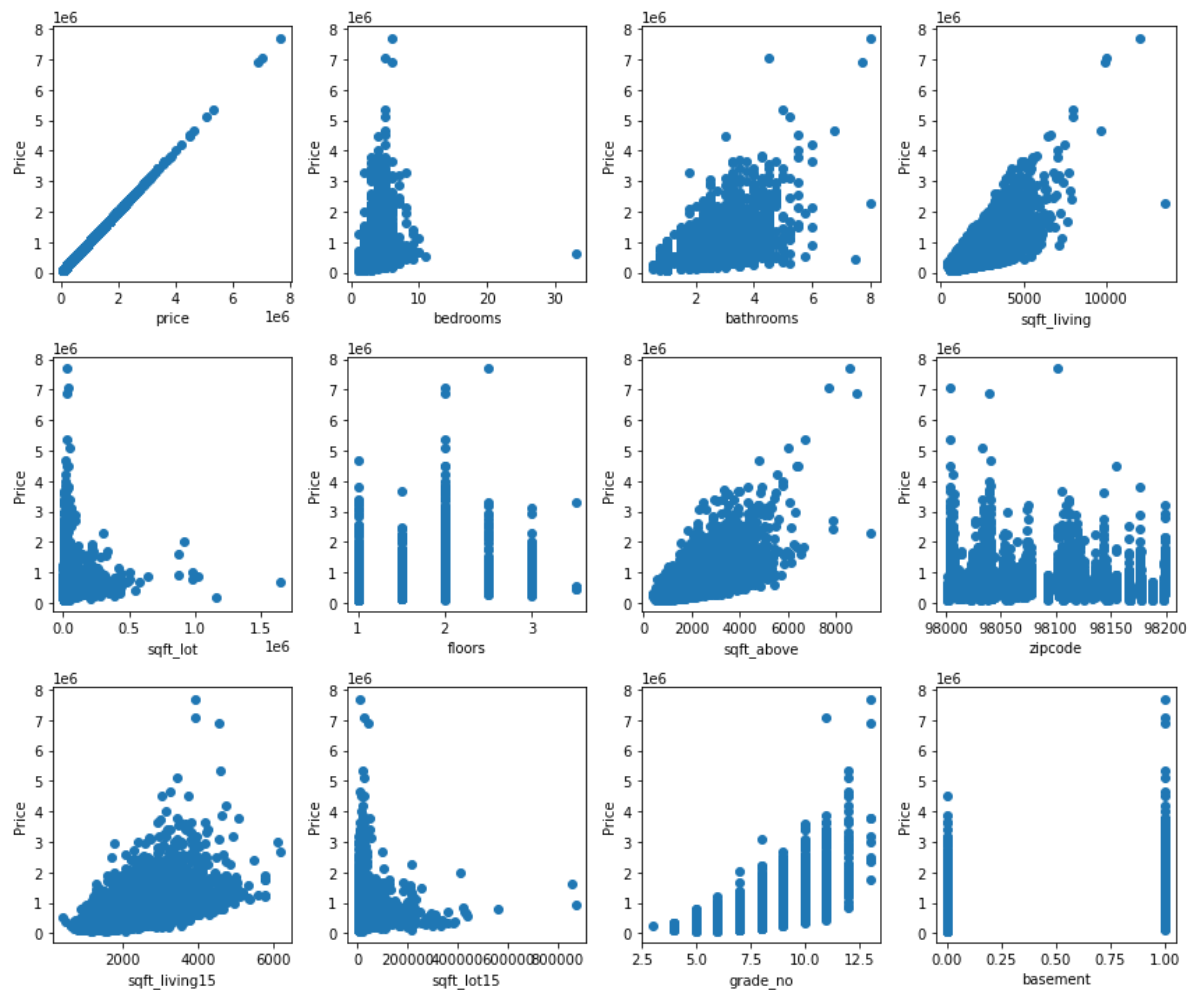
In [47]: # Plot scatter plots against "price"
X = df
y = df["price"]

fig, axes = plt.subplots(nrows=3, ncols=4, figsize=(12, 10))
flatten_axes = axes.flatten()

for i, column in enumerate(X.columns):
    flatten_axes[i].scatter(X[column], y)
    flatten_axes[i].set_xlabel(column)
    flatten_axes[i].set_ylabel("Price")

plt.tight_layout()
plt.show()

```



```
In [48]: continuous = ['price', 'sqft_living', 'sqft_lot', 'sqft_living15', 'sqft_lot15']
df_no_outlier = outliers(continuous, df)
df_no_outlier.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 14582 entries, 1 to 21596
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  ---
0   price            14582 non-null  float64
1   bedrooms         14582 non-null  int64
2   bathrooms        14582 non-null  float64
3   sqft_living      14582 non-null  int64
4   sqft_lot         14582 non-null  int64
5   floors           14582 non-null  float64
6   sqft_above       14582 non-null  int64
7   zipcode          14582 non-null  int64
8   sqft_living15    14582 non-null  int64
9   sqft_lot15       14582 non-null  int64
10  grade_no         14582 non-null  int64
11  basement         14582 non-null  int32
dtypes: float64(3), int32(1), int64(8)
memory usage: 1.4 MB
```

Iteration 2

In this iteration, we tried to remove outliers from our data to see the impact on our model's performance.

In [49]: *# Fitting our model without outliers*

```
run_model(df_no_outlier)
```

The mean r^2 for a KFold test with 10 splits is 0.5023419026527095

The mean RMSE for a KFold test with 10 splits is 173691.9672560325

OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.50
Model:                  OLS      Adj. R-squared:            0.50
Method:                 Least Squares    F-statistic:          134
Date:                   Thu, 01 Jun 2023    Prob (F-statistic):      0.0
Time:                   12:27:04    Log-Likelihood:         -1.9661e+0
No. Observations:       14582    AIC:                    3.932e+0
Df Residuals:           14570    BIC:                    3.933e+0
Df Model:               11
Covariance Type:        nonrobust
=====
```


```
=====
=====
coef      std err          t      P>|t|      [0.025      0.
975]
-----
Intercept    -5.059e+07    2.85e+06    -17.728    0.000    -5.62e+07    -4.5
e+07
bedrooms     -1.882e+04    2190.043     -8.595    0.000    -2.31e+04    -1.45
e+04
bathrooms    -2.732e+04    3422.967     -7.982    0.000    -3.4e+04     -2.06
e+04
sqft_living   155.8640         6.919     22.528    0.000     142.302      16
9.426
sqft_lot       0.3337         0.230      1.448    0.148      -0.118
0.786
floors        7402.9042    3829.544      1.933    0.053    -103.488     1.49
e+04
sqft_above    -47.7686         7.463     -6.400    0.000     -62.398      -3
3.139
zipcode       511.4977        29.081     17.589    0.000     454.495      56
8.500
sqft_living15  64.7177         3.987     16.232    0.000      56.903       7
2.533
sqft_lot15    -2.3417         0.326     -7.175    0.000     -2.981       -
1.702
grade_no      9.022e+04    2259.478     39.930    0.000     8.58e+04     9.46
e+04
basement      2.111e+04    5466.456      3.861    0.000     1.04e+04     3.18
e+04
=====
```


```
=====
Omnibus:          2782.963    Durbin-Watson:          1.96
5
```


Prob(Omnibus):	0.000	Jarque-Bera (JB):	7255.79
9			
Skew:	1.041	Prob(JB):	0.0
0			
Kurtosis:	5.759	Cond. No.	1.97e+08
8			
=====			
=			

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The condition number is large, 1.97e+08. This might indicate that there are strong multicollinearity or other numerical problems.





Interpretation of results

1. The model is generally statistically significant with an F-statistic p_value of 0.0 at a significance level of 0.05
2. The R-squared value is 0.504, indicating that approximately 50.4% of the variation in the price can be explained by the model. This value indicates a drop from the previous model.
3. The plot to test for homoscedasticity reveals that the residuals are becoming homoscedastic because they are converging and appear to be having an equal variance. So this assumption is satisfied.
4. The QQ-plot is used to test for normality of residuals. In this case, the residuals appear to be somewhat normal but there is still presence of skewness/heavy-tails/outliers.

```
In [50]: # Displaying the DataFrame
df_no_outlier
```

```
Out[50]:
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	sqft_above	zipcode	sqft_liv
1	538000.0	3	2.25	2570	7242	2.0	2170	98125	
3	604000.0	4	3.00	1960	5000	1.0	1050	98136	
4	510000.0	3	2.00	1680	8080	1.0	1680	98074	
6	257500.0	3	2.25	1715	6819	2.0	1715	98003	
8	229500.0	3	1.00	1780	7470	1.0	1050	98146	
...
21591	475000.0	3	2.50	1310	1294	2.0	1180	98116	
21592	360000.0	3	2.50	1530	1131	3.0	1530	98103	
21593	400000.0	4	2.50	2310	5813	2.0	2310	98146	
21594	402101.0	2	0.75	1020	1350	2.0	1020	98144	
21596	325000.0	2	0.75	1020	1076	2.0	1020	98144	

14582 rows × 12 columns

```
In [51]: df_no_outlier.info()
```


```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 14582 entries, 1 to 21596
Data columns (total 12 columns):
#   Column              Non-Null Count  Dtype
---  ---
0   price               14582 non-null  float64
1   bedrooms            14582 non-null  int64
2   bathrooms           14582 non-null  float64
3   sqft_living         14582 non-null  int64
4   sqft_lot            14582 non-null  int64
5   floors              14582 non-null  float64
6   sqft_above          14582 non-null  int64
7   zipcode             14582 non-null  int64
8   sqft_living15       14582 non-null  int64
9   sqft_lot15          14582 non-null  int64
10  grade_no            14582 non-null  int64
11  basement            14582 non-null  int32
dtypes: float64(3), int32(1), int64(8)
memory usage: 1.4 MB
```

```
In [52]: # Dropping unnecessary columns
df_no_outlier.drop(columns=["sqft_above", "grade_no"], inplace=True)
```

```
In [53]: df_no_outlier.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 14582 entries, 1 to 21596
Data columns (total 10 columns):
 #   Column          Non-Null Count  Dtype  
---  -
 0   price           14582 non-null  float64
 1   bedrooms        14582 non-null  int64  
 2   bathrooms        14582 non-null  float64
 3   sqft_living      14582 non-null  int64  
 4   sqft_lot         14582 non-null  int64  
 5   floors           14582 non-null  float64
 6   zipcode          14582 non-null  int64  
 7   sqft_living15    14582 non-null  int64  
 8   sqft_lot15       14582 non-null  int64  
 9   basement         14582 non-null  int32  
dtypes: float64(3), int32(1), int64(6)
memory usage: 1.2 MB
```

```
In [54]: # Displaying the heatmap
heatmap(df_no_outlier)
```



Iteration 3

In this iteration, we perform some normalization and log-transformations. This will help to mitigate the presence of outliers in our dataset and hence make the dataset more robust, and also improving the linearity between the target variable(price) and the features.

Normalization and Log_transformation

In [55]: *# Displaying the DataFrame*
df_no_outlier

Out[55]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	zipcode	sqft_living15	sqft_l
1	538000.0	3	2.25	2570	7242	2.0	98125	1690	
3	604000.0	4	3.00	1960	5000	1.0	98136	1360	
4	510000.0	3	2.00	1680	8080	1.0	98074	1800	
6	257500.0	3	2.25	1715	6819	2.0	98003	2238	
8	229500.0	3	1.00	1780	7470	1.0	98146	1780	
...
21591	475000.0	3	2.50	1310	1294	2.0	98116	1330	
21592	360000.0	3	2.50	1530	1131	3.0	98103	1530	
21593	400000.0	4	2.50	2310	5813	2.0	98146	1830	
21594	402101.0	2	0.75	1020	1350	2.0	98144	1020	
21596	325000.0	2	0.75	1020	1076	2.0	98144	1020	

14582 rows × 10 columns

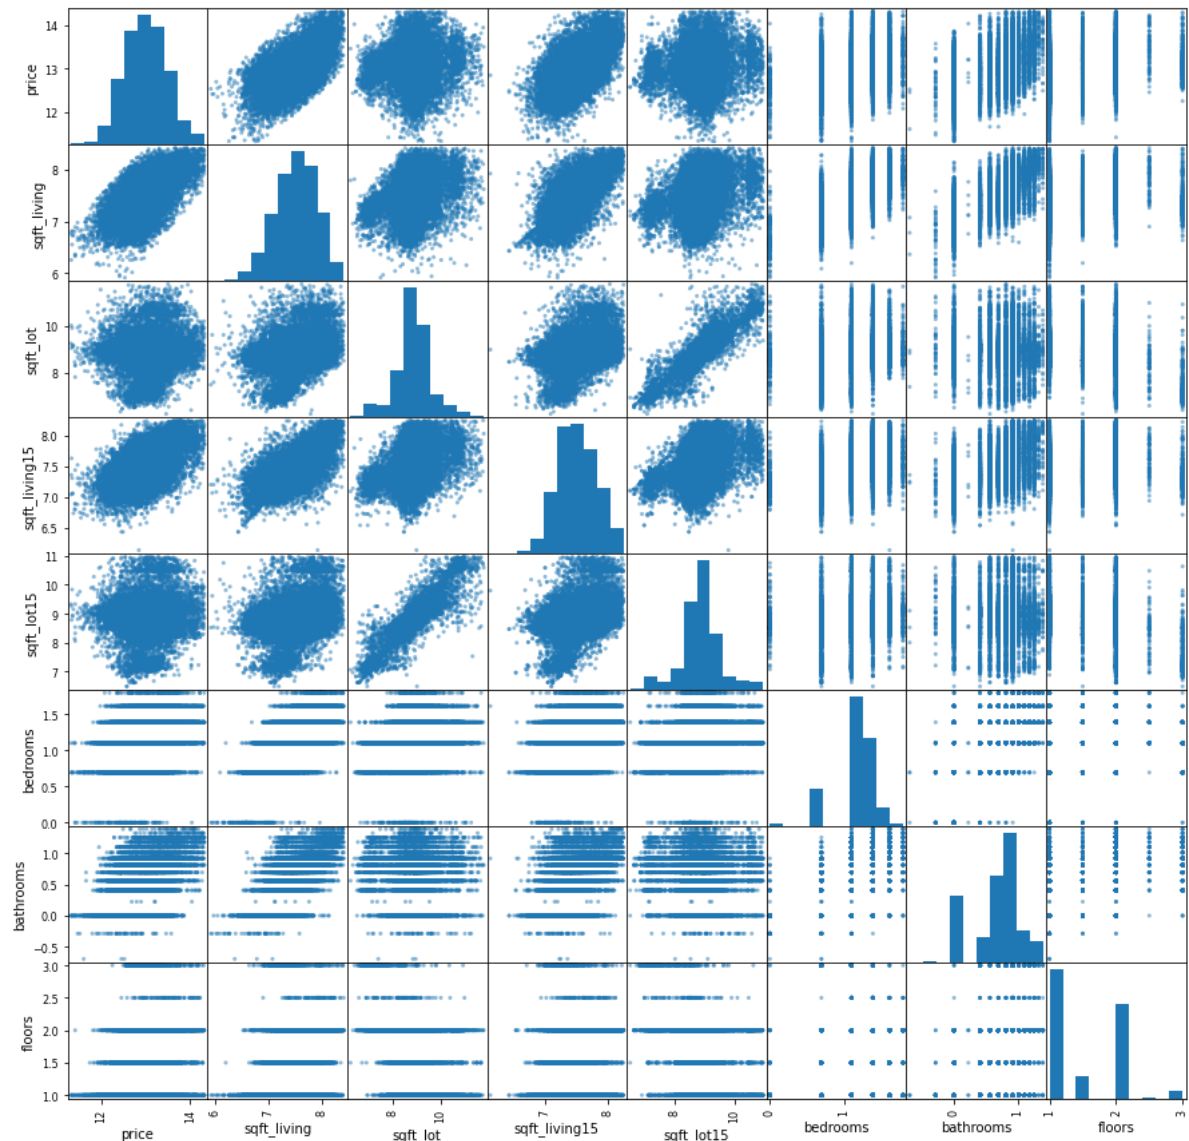
In [56]: *# Checking the correlations in descending order*
df_no_outlier.corr()["price"].sort_values(ascending=False)

Out[56]:

price	1.000000
sqft_living	0.627050
sqft_living15	0.550582
bathrooms	0.446199
bedrooms	0.297462
floors	0.256793
basement	0.178915
sqft_lot15	0.093464
sqft_lot	0.091582
zipcode	-0.005953

Name: price, dtype: float64

```
In [57]: # Performing Log transformations using our defined function
normalize = ['price', 'sqft_living', 'sqft_lot', 'sqft_living15', 'sqft_lot15']
df_log = log_transform(normalize, df_no_outlier)
pd.plotting.scatter_matrix(df_log[continuous], figsize=(15, 15));
```



```
In [58]: # Using our `df_log` we fit our model using our defined function  
run_model(df_log)
```


The mean r^2 for a KFold test with 10 splits is 0.45540890401512824

The mean RMSE for a KFold test with 10 splits is 0.348485819455707

OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.45
Model:                  OLS      Adj. R-squared:            0.45
Method:                 Least Squares    F-statistic:        136
Date:                   Thu, 01 Jun 2023    Prob (F-statistic):    0.0
Time:                   12:27:27    Log-Likelihood:       -5309.
No. Observations:      14582    AIC:                  1.064e+0
Df Residuals:          14572    BIC:                  1.072e+0
Df Model:               9
Covariance Type:        nonrobust
=====
```

```
=====
=====
coef      std err          t      P>|t|      [0.025      0.
975]
-----
Intercept    -100.5845      5.805    -17.326      0.000    -111.964    -8
bedrooms     -0.1838      0.014    -13.162      0.000     -0.211     -
bathrooms    -0.0522      0.013     -4.114      0.000     -0.077     -
sqft_living   0.6269      0.016    39.088      0.000      0.595
sqft_lot     -0.0268      0.010     -2.766      0.006     -0.046     -
floors        0.0417      0.008      5.237      0.000      0.026
zipcode       0.0011     5.9e-05    18.355      0.000      0.001
sqft_living15  0.4777      0.015    32.520      0.000      0.449
sqft_lot15   -0.0669      0.011     -6.311      0.000     -0.088     -
basement      0.0685      0.007      9.315      0.000      0.054
=====
```

```
=====
Omnibus:          52.603    Durbin-Watson:          1.99
Prob(Omnibus):    0.000    Jarque-Bera (JB):       40.05
Skew:             -0.009    Prob(JB):               2.01e-0
=====
```

Kurtosis:

2.744

Cond. No.

1.97e+08

8



=====

=

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.


[2] The condition number is large, 1.97e+08. This might indicate that there are strong multicollinearity or other numerical problems.



Interpretation of results

1. The model is generally statistically significant with an F-statistic p_value of 0.0 at a significance level of 0.05
2. The R-squared value is 0.457, indicating that approximately 45.7% of the variation in the price can be explained by the model. This value indicates a drop from the previous model.
3. The plot to test for homoscedasticity reveals that the residuals are now homoscedastic because they are converging and appear to be having an equal variance. So this assumption remains satisfied.
4. The QQ-plot is used to test for normality of residuals. In this case, the residuals appear to be almost perfectly normal as they are following along the line almost neatly.

```
In [59]: # Displaying the heatmap
heatmap(df_log)
```



Iteration 4 (Final Model)

One hot encode Zipcode

```
In [60]: # Define a function to perform feature scaling
def scale(feature):
    return (feature-feature.min())/(feature.max()-feature.min())

# OneHotEncoding zipcode
df_scale = scale_ohe('zipcode', df_log)
```

In [61]: `df_scale.info()`

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 14582 entries, 1 to 21596
Data columns (total 78 columns):
 #   Column                Non-Null Count  Dtype
---  -
 0   price                 14582 non-null  float64
 1   bedrooms              14582 non-null  float64
 2   bathrooms             14582 non-null  float64
 3   sqft_living           14582 non-null  float64
 4   sqft_lot              14582 non-null  float64
 5   floors                14582 non-null  float64
 6   sqft_living15         14582 non-null  float64
 7   sqft_lot15            14582 non-null  float64
 8   basement              14582 non-null  float64
 9   zipcode_98002         14582 non-null  uint8
10   zipcode_98003         14582 non-null  uint8
11   zipcode_98004         14582 non-null  uint8
12   zipcode_98005         14582 non-null  uint8
13   zipcode_98006         14582 non-null  uint8
14   zipcode_98007         14582 non-null  uint8
15   zipcode_98008         14582 non-null  uint8
16   zipcode_98010         14582 non-null  uint8
17   zipcode_98011         14582 non-null  uint8
18   zipcode_98014         14582 non-null  uint8
19   zipcode_98019         14582 non-null  uint8
20   zipcode_98022         14582 non-null  uint8
21   zipcode_98023         14582 non-null  uint8
22   zipcode_98024         14582 non-null  uint8
23   zipcode_98027         14582 non-null  uint8
24   zipcode_98028         14582 non-null  uint8
25   zipcode_98029         14582 non-null  uint8
26   zipcode_98030         14582 non-null  uint8
27   zipcode_98031         14582 non-null  uint8
28   zipcode_98032         14582 non-null  uint8
29   zipcode_98033         14582 non-null  uint8
30   zipcode_98034         14582 non-null  uint8
31   zipcode_98038         14582 non-null  uint8
32   zipcode_98039         14582 non-null  uint8
33   zipcode_98040         14582 non-null  uint8
34   zipcode_98042         14582 non-null  uint8
35   zipcode_98045         14582 non-null  uint8
36   zipcode_98052         14582 non-null  uint8
37   zipcode_98053         14582 non-null  uint8
38   zipcode_98055         14582 non-null  uint8
39   zipcode_98056         14582 non-null  uint8
40   zipcode_98058         14582 non-null  uint8
41   zipcode_98059         14582 non-null  uint8
42   zipcode_98065         14582 non-null  uint8
43   zipcode_98070         14582 non-null  uint8
44   zipcode_98072         14582 non-null  uint8
45   zipcode_98074         14582 non-null  uint8
46   zipcode_98075         14582 non-null  uint8
47   zipcode_98077         14582 non-null  uint8
48   zipcode_98092         14582 non-null  uint8
49   zipcode_98102         14582 non-null  uint8
50   zipcode_98103         14582 non-null  uint8
51   zipcode_98105         14582 non-null  uint8

```

```
52  zipcode_98106  14582 non-null  uint8
53  zipcode_98107  14582 non-null  uint8
54  zipcode_98108  14582 non-null  uint8
55  zipcode_98109  14582 non-null  uint8
56  zipcode_98112  14582 non-null  uint8
57  zipcode_98115  14582 non-null  uint8
58  zipcode_98116  14582 non-null  uint8
59  zipcode_98117  14582 non-null  uint8
60  zipcode_98118  14582 non-null  uint8
61  zipcode_98119  14582 non-null  uint8
62  zipcode_98122  14582 non-null  uint8
63  zipcode_98125  14582 non-null  uint8
64  zipcode_98126  14582 non-null  uint8
65  zipcode_98133  14582 non-null  uint8
66  zipcode_98136  14582 non-null  uint8
67  zipcode_98144  14582 non-null  uint8
68  zipcode_98146  14582 non-null  uint8
69  zipcode_98148  14582 non-null  uint8
70  zipcode_98155  14582 non-null  uint8
71  zipcode_98166  14582 non-null  uint8
72  zipcode_98168  14582 non-null  uint8
73  zipcode_98177  14582 non-null  uint8
74  zipcode_98178  14582 non-null  uint8
75  zipcode_98188  14582 non-null  uint8
76  zipcode_98198  14582 non-null  uint8
77  zipcode_98199  14582 non-null  uint8
dtypes: float64(9), uint8(69)
memory usage: 2.1 MB
```



```
In [62]: # Using our `df_scale` we fit our model using our defined function  
run_model(df_scale)
```

The mean r^2 for a KFold test with 10 splits is 0.8313639887029263

The mean RMSE for a KFold test with 10 splits is 0.06464512379839615

OLS Regression Results

```
=====
Dep. Variable:          price    R-squared:                0.83
Model:                  OLS      Adj. R-squared:            0.83
Method:                 Least Squares    F-statistic:        942.
Date:                   Thu, 01 Jun 2023    Prob (F-statistic):    0.0
Time:                   12:27:36    Log-Likelihood:        1933
No. Observations:      14582    AIC:                   -3.851e+0
Df Residuals:          14504    BIC:                   -3.792e+0
Df Model:               77
Covariance Type:        nonrobust
=====
```

```
=====
=====
coef      std err          t      P>|t|      [0.025      0.
975]
-----
----
Intercept    -0.0694      0.006    -12.052    0.000    -0.081    -
0.058
bedrooms     -0.0545      0.005    -11.575    0.000    -0.064    -
0.045
bathrooms     0.0552      0.005     11.066    0.000     0.045
0.065
sqft_living   0.4423      0.008    58.880    0.000     0.428
0.457
sqft_lot      0.1140      0.010     11.398    0.000     0.094
0.134
floors        0.0086      0.003     2.830    0.005     0.003
0.015
sqft_living15 0.1799      0.006    29.385    0.000     0.168
0.192
sqft_lot15    0.0042      0.009     0.459    0.646    -0.014
0.022
basement     -0.0154      0.001    -10.806    0.000    -0.018    -
0.013
zipcode_98002 0.0119      0.007     1.763    0.078    -0.001
0.025
zipcode_98003 0.0189      0.006     3.136    0.002     0.007
0.031
zipcode_98004 0.3834      0.007    58.608    0.000     0.371
0.396
zipcode_98005 0.2705      0.007    37.036    0.000     0.256
0.285
zipcode_98006 0.2437      0.006    43.357    0.000     0.233
0.255
=====
```

zipcode_98007	0.2402	0.008	31.230	0.000	0.225
0.255					
zipcode_98008	0.2428	0.006	39.884	0.000	0.231
0.255					
zipcode_98010	0.0875	0.010	9.029	0.000	0.069
0.107					
zipcode_98011	0.1522	0.007	22.060	0.000	0.139
0.166					
zipcode_98014	0.1025	0.009	10.797	0.000	0.084
0.121					
zipcode_98019	0.0964	0.007	13.001	0.000	0.082
0.111					
zipcode_98022	0.0250	0.007	3.599	0.000	0.011
0.039					
zipcode_98023	0.0075	0.005	1.410	0.158	-0.003
0.018					
zipcode_98024	0.1459	0.013	11.346	0.000	0.121
0.171					
zipcode_98027	0.1944	0.006	33.116	0.000	0.183
0.206					
zipcode_98028	0.1465	0.006	24.215	0.000	0.135
0.158					
zipcode_98029	0.2203	0.006	37.343	0.000	0.209
0.232					
zipcode_98030	0.0226	0.006	3.588	0.000	0.010
0.035					
zipcode_98031	0.0305	0.006	4.971	0.000	0.018
0.042					
zipcode_98032	0.0124	0.008	1.613	0.107	-0.003
0.028					
zipcode_98033	0.2797	0.006	50.254	0.000	0.269
0.291					
zipcode_98034	0.1965	0.005	37.396	0.000	0.186
0.207					
zipcode_98038	0.0516	0.005	9.826	0.000	0.041
0.062					
zipcode_98039	0.4344	0.019	22.777	0.000	0.397
0.472					
zipcode_98040	0.3286	0.006	50.565	0.000	0.316
0.341					
zipcode_98042	0.0289	0.005	5.512	0.000	0.019
0.039					
zipcode_98045	0.1170	0.007	16.992	0.000	0.103
0.130					
zipcode_98052	0.2313	0.005	44.223	0.000	0.221
0.242					
zipcode_98053	0.2062	0.006	35.192	0.000	0.195
0.218					
zipcode_98055	0.0582	0.006	9.434	0.000	0.046
0.070					
zipcode_98056	0.1253	0.006	22.611	0.000	0.114
0.136					
zipcode_98058	0.0624	0.005	11.394	0.000	0.052
0.073					
zipcode_98059	0.1194	0.006	21.593	0.000	0.109
0.130					
zipcode_98065	0.1376	0.006	22.090	0.000	0.125


0.150					
zipcode_98070	0.1627	0.010	16.443	0.000	0.143
0.182					
zipcode_98072	0.1648	0.006	26.083	0.000	0.152
0.177					
zipcode_98074	0.2082	0.006	36.642	0.000	0.197
0.219					
zipcode_98075	0.2095	0.006	35.390	0.000	0.198
0.221					
zipcode_98077	0.1607	0.007	22.127	0.000	0.146
0.175					
zipcode_98092	0.0135	0.006	2.263	0.024	0.002
0.025					
zipcode_98102	0.3627	0.010	38.009	0.000	0.344
0.381					
zipcode_98103	0.3123	0.005	58.738	0.000	0.302
0.323					
zipcode_98105	0.3547	0.007	53.284	0.000	0.342
0.368					
zipcode_98106	0.1464	0.006	24.548	0.000	0.135
0.158					
zipcode_98107	0.3255	0.006	52.372	0.000	0.313
0.338					
zipcode_98108	0.1429	0.007	20.479	0.000	0.129
0.157					
zipcode_98109	0.3697	0.009	42.620	0.000	0.353
0.387					
zipcode_98112	0.3825	0.007	58.276	0.000	0.370
0.395					
zipcode_98115	0.3066	0.005	58.043	0.000	0.296
0.317					
zipcode_98116	0.3017	0.006	51.243	0.000	0.290
0.313					
zipcode_98117	0.3097	0.005	58.388	0.000	0.299
0.320					
zipcode_98118	0.1845	0.005	34.233	0.000	0.174
0.195					
zipcode_98119	0.3771	0.007	52.277	0.000	0.363
0.391					
zipcode_98122	0.3087	0.006	49.448	0.000	0.296
0.321					
zipcode_98125	0.2109	0.006	37.772	0.000	0.200
0.222					
zipcode_98126	0.2306	0.006	39.374	0.000	0.219
0.242					
zipcode_98133	0.1740	0.005	32.208	0.000	0.163
0.185					
zipcode_98136	0.2746	0.006	43.947	0.000	0.262
0.287					
zipcode_98144	0.2600	0.006	43.693	0.000	0.248
0.272					
zipcode_98146	0.1162	0.006	19.196	0.000	0.104
0.128					
zipcode_98148	0.0495	0.011	4.554	0.000	0.028
0.071					
zipcode_98155	0.1584	0.006	28.685	0.000	0.148
0.169					


zipcode_98166	0.1393	0.006	21.996	0.000	0.127
0.152					
zipcode_98168	0.0378	0.006	6.052	0.000	0.026
0.050					
zipcode_98177	0.2331	0.006	36.573	0.000	0.221
0.246					
zipcode_98178	0.0762	0.006	12.109	0.000	0.064
0.089					
zipcode_98188	0.0381	0.008	4.926	0.000	0.023
0.053					
zipcode_98198	0.0475	0.006	7.728	0.000	0.035
0.060					
zipcode_98199	0.3261	0.006	53.011	0.000	0.314
0.338					

```
=====
=
Omnibus:                981.961    Durbin-Watson:                1.99
0
Prob(Omnibus):           0.000    Jarque-Bera (JB):            4604.77
8
Skew:                    -0.124    Prob(JB):                     0.0
0
Kurtosis:                5.742    Cond. No.                     12
1.
=====
=
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.





```
In [63]: def get_coefficients_categorical(scaled_coefs, features):
    for i, feat in enumerate(features):
        maximum = df_log['price'].max()
        minimum = df_log['price'].min()
        unscale = abs(scaled_coefs[i])*(maximum-minimum)+minimum
        unlog = math.exp(unscale)
        if scaled_coefs[i] >= 0:
            print('Coefficient for {} is ${}'.format(feat, unlog))
        else:
            print('Coefficient for {} is ${}'.format(feat, unlog*-1))
```

```
In [64]: def get_coefficients_continuous(scaled_coefs, features):
    for i, feat in enumerate(features):
        maximum = df_log['price'].max()
        minimum = df_log['price'].min()
        range_feat = df_no_outlier[feat].max() - df_no_outlier[feat].min()
        unscale = abs(scaled_coefs[i])*(maximum-minimum)+minimum
        unlog = math.exp(unscale)

        slope_actual = unlog/range_feat

        if scaled_coefs[i] >= 0:
            print('Coefficient for {} is ${}'.format(feat, slope_actual))
        else:
            print('Coefficient for {} is ${}'.format(feat, slope_actual*-1))
```

```
In [79]: categorical_coef = [0.2428, 0.2082, 0.3097, 0.2600, 0.2331]
categorical_features = ['zipcode_98008', 'zipcode_98074', 'zipcode_98117', 'zip

continuous_coef = [0.4423, 0.1799, -0.0545, 0.0086, 0.0552, 0.1140]
continuous_features = ['sqft_living', 'sqft_living15', 'bedrooms', 'floors', 'l

get_coefficients_categorical(categorical_coef, categorical_features)
get_coefficients_continuous(continuous_coef, continuous_features)
```

```
Coefficient for zipcode_98008 is $169959.30663666
Coefficient for zipcode_98074 is $153192.66292287616
Coefficient for zipcode_98117 is $207759.30309087687
Coefficient for zipcode_98144 is $178964.988254935
Coefficient for zipcode_98177 is $165081.8589739885
Coefficient for sqft_living is $123487.74911877913
Coefficient for sqft_living15 is $66394.5581188671
Coefficient for bedrooms is $-53899.42895874723
Coefficient for floors is $42072.219101705305
Coefficient for bathrooms is $46540.35864547536
Coefficient for sqft_lot is $20891.42559555274
```

Train Test Split

```
In [65]: # Getting a copy of our df
df_tts = df.copy()
x = df_tts.drop('price', axis=1)
y = df_tts['price']
```

Split original data into training data (80%) and testing data (20%).

```
In [66]: # Split the data
x_train, x_test, y_train, y_test = train_test_split(x, y, test_size=0.20)
```

Concat x with y to remove outliers

```
In [67]: # Concat x with y to remove outliers
train = pd.concat([x_train, y_train], axis=1)
test = pd.concat([x_test, y_test], axis=1)
len(train)
```

Out[67]: 12540

Remove outliers separately

```
In [68]: # Remove outliers separately
train1 = outliers(continuous, train)
test1 = outliers(continuous, test)
len(train1)
```

Out[68]: 11646

Log transform train and test splits

```
In [69]: # Log transform train and test splits
train2 = log_transform(normalize, train1)
test2 = log_transform(normalize, test1)
```

Scale and OHE training and testing data separately

```
In [70]: # Scale and OHE training and testing data separately
train_preprocessed = scale_ohe('zipcode', train2)

test_preprocessed = scale_ohe('zipcode', test2)
```

Drop features determined by our final model

```
In [71]: # Drop features determined by our final model
train_preprocessed.drop(['sqft_lot15', 'zipcode_98002', 'zipcode_98023', 'zipcode_98044'], axis=1)
test_preprocessed.drop(['sqft_lot15', 'zipcode_98002', 'zipcode_98023', 'zipcode_98044'], axis=1)
```

Apply interactions determined by our final model

```
In [72]: # Apply interactions determined by our final model
train_preprocessed['sqft_living*floors'] = train_preprocessed['sqft_living']*train_preprocessed['floors']
test_preprocessed['sqft_living*floors'] = test_preprocessed['sqft_living']*test_preprocessed['floors']
```

Check to see that the training and testing sets are split correctly

```
In [80]: # Check to see that the training and testing sets are split correctly
x_train_preprocessed = train_preprocessed.drop('price', axis=1)
y_train_preprocessed = train_preprocessed['price']

x_test_preprocessed = test_preprocessed.drop('price', axis=1)
y_test_preprocessed = test_preprocessed['price']

print(len(x_train_preprocessed), len(x_test_preprocessed), len(y_train_preprocessed), len(y_test_preprocessed))

11646 2902 11646 2902
```

Run testing data through training model

```
In [84]: # Run testing data through training model
linreg = LinearRegression()
linreg.fit(x_train_preprocessed, y_train_preprocessed)
y_hat_test = linreg.predict(x_test_preprocessed)

test_rmse = mean_squared_error(y_test_preprocessed, y_hat_test, squared=False)
test_rmse
```

Out[84]: 0.06463831997659172


```
In [83]: # Calculate evaluation metrics on the original scale
y_pred_original = np.exp(y_hat_test) # Transform predicted values back to the
y_test_original = np.exp(y_test_preprocessed) # Transform actual values back to the original scale

rmse_original = mean_squared_error(y_test_original, y_pred_original, squared=False)

print("RMSE in original scale:", rmse_original)
```

RMSE in original scale: 0.11351287323753316

CONCLUSIONS

Interpretation of results from the Final Model

1. The model is generally statistically significant with an F-statistic p_value of 0.0 at a significance level of 0.05
2. The R-squared value is 0.833, indicating that approximately 83.3% of the variation in the price can be explained by the model. This value indicates a great improvement from the previous model.
3. Also, of great importance to note is that the mean RMSE is approximately 0.06465. Then the RMSE in original scale is 0.1135. This means that our model is off by about 0.1135 when making an average prediction, indicating that it is a good model.
4. These coefficients represent the expected change in the price for a one-unit change in the corresponding predictor variable, assuming other variables are held constant.
 - ZIPCODE--is a strong predictor of a homes value, the saying "Location, Location, Location" holds true, as even in a similar area the location plays a huge factor in the value of a home.

Based on the coefficients of different localities, moving from zip code 98002 to 98039 shows that the prices changes by USD 228,087 and USD 298,174 respectively, as compared to our reference categorical variable which is zipcode 98001. This is a clear indication that locality of the house has high influence on the price.

- Coefficient for `sqft_living` is \$123487.74911877913
 - For a one-unit increase in square-foot living area, we see an associated increase in around \$123487.74 in selling price of the houses.
- Coefficient for `sqft_living15` is \$66394.5581188671
 - For a one-unit increase in square-foot living area15, we see an associated increase in around \$66394.55 in selling price of the houses.
- Coefficient for `floors` is \$42072.219101705305
 - For a one-unit increase in number of floors of the house, we see an associated increase in around \$42072.21 in selling price of the houses.
- Coefficient for `bathrooms` is \$46540.35864547536
 - For a one-unit increase in the number of bathrooms, we see an associated increase in around \$46540.35 in selling price of the houses.

- Coefficient for `sqft_lot` is \$20891.42559555274
 - For a one-unit increase in square-foot of the lot area, we see an associated increase in around \$20891.42 in selling price of the houses.
 - Coefficient for `bedrooms` is \$-53899.42895874723
 - For a one-unit increase in the number of bedrooms, we see an associated decrease in around \$53899.42 in selling price of the houses. This particular finding caught our attention as this is not the case in the real world, whereby as you increase the number of bedrooms in a house, the price of the house tends to increase too.
5. The plot to test for homoscedasticity reveals that the residuals are now homoscedastic because they are converging and appear to be having an equal variance. So this assumption remains satisfied.
 6. The QQ-plot is used to test for normality of residuals. In this case, the residuals appear to be almost normal as they are following along the line almost neatly, except for the ends where it indicates there could be some skewness in the data.

RECOMMENDATIONS

1. The real estate agency should explore properties that occupy a large square foot of the lot area since, for a one-unit increase in square-foot of the lot area, we see an associated increase in around \$ 20891.42 in selling price of the houses.
2. The real estate agency should explore properties that have more bathrooms since, for a one-unit increase in the number of bathrooms, we see an associated increase in around \$ 46540.35 in selling price of the houses.
3. The real estate agency should explore properties that occupy a large square foot of living area since, for a one-unit increase in square-foot living area, we see an associated increase in around \$ 123487.74 in selling price of the houses.
4. The real estate agency should explore properties with more floors since, for a one-unit increase in number of floors of the house, we see an associated increase in around \$ 42072.21 in selling price of the houses.

NEXT STEPS

1. More research is required to have a more integrated and informative dataset for finding more factors that influence the price. Also, use of more complex and robust regression models that will help to deal with the outliers.
2. Using datasets from other counties to be able to better advice our customers from comparing the dataset results.
3. It is also important for the agency to continuously evaluate the effectiveness of the strategies they implement and make adjustments as necessary. This could involve tracking

metrics like, this model, social media engagement/reviews, and lead generation to assess the impact of their efforts and identify areas for improvement.